

## ▼ CS445: Computational Photography

### Programming Project 4: Image-Based Lighting

#### ▼ Recovering HDR Radiance Maps

Load libraries and data

```
1 # jupyter extension that allows reloading functions from imports without clearing
2 %load_ext autoreload
3 %autoreload 2
```

The autoreload extension is already loaded. To reload it, use:  
  %reload\_ext autoreload

```
1 from google.colab import drive
2 drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call dr

```
1 # System imports
2 from os import path
3 import math
4
5 # Third-Party Imports
6 import cv2
7 import matplotlib.pyplot as plt
8 import numpy as np
9 from scipy.interpolate import griddata
10 import random
11
12 # modify to where you store your project data including utils
13 datadir = "/content/drive/My Drive/Semesters/Semester 7/CS445/MPs/MP4/proj4/"
14
15 utilfn = datadir + "utils"
16 !cp -r "$utilfn" .
17 samplesfn = datadir + "samples"
18 !cp -r "$samplesfn" .
19 inputsfn = datadir + "inputs"
20 !cp -r "$inputsfn" .
21
22 # can change this to your output directory of choice
23 !mkdir "images"
24 !mkdir "images/outputs"
```

```

25
26 # import starter code
27 import utils
28 from utils.io import read_image, write_image, read_hdr_image, write_hdr_image
29 from utils.display import display_images_linear_rescale, rescale_images_linear
30 from utils.hdr_helpers import gsolve
31 from utils.hdr_helpers import get_equirectangular_image
32 from utils.bilateral_filter import bilateral_filter
33
34

mkdir: cannot create directory 'images': File exists
mkdir: cannot create directory 'images/outputs': File exists

```

## ▼ Reading LDR images

You can use the provided samples or your own images. You get more points for using your own images, but it might help to get things working first with the provided samples.

```

1 # TODO: Replace this with your path and files
2
3 imdir = 'inputs'
4
5 imfns = ['200.jpg', '400.jpg', '800.jpg', '1600.jpg', '3200.jpg', '6400.jpg']
6 exposure_times = [1/200.0, 1/400.0, 1/800.0, 1/1600.0, 1/3200.0, 1/6400.0]
7 background_image_file = imdir + '/' + 'baseline.jpg'
8
9 # Use this for other light map (bells & whistles)
10 # imfns = ['8.jpg', '10.jpg', '40.jpg', '160.jpg']
11 # exposure_times = [1/8.0, 1/10.0, 1/40.0, 1/160.0]
12 # background_image_file = imdir + '/' + 'baseline3.jpg'
13
14 ldr_images = []
15 for f in np.arange(len(imfns)):
16     im = read_image(imdir + '/' + imfns[f])
17     if f==0:
18         imsize = int((im.shape[0] + im.shape[1])/2) # set width/height of ball images
19         ldr_images = np.zeros((len(imfns), imsize, imsize, 3))
20     ldr_images[f] = cv2.resize(im, (imsize, imsize))
21
22 background_image = read_image(background_image_file)
23
24

```

## ▼ Naive LDR merging

Compute the HDR image as average of irradiance estimates from LDR images

```

1 def make_hdr_naive(ldr_images: np.ndarray, exposures: list) -> (np.ndarray, np.nda
2     """
3         Makes HDR image using multiple LDR images, and its corresponding exposure val
4
5         The steps to implement:
6             1) Divide each image by its exposure time.
7                 - This will rescale images as if it has been exposed for 1 second.
8
9
10
11             2) Return average of above images
12
13
14
15             For further explanation, please refer to problem page for how to do it.
16
17             Args:
18                 ldr_images(np.ndarray): N x H x W x 3 shaped numpy array representing
19                     N ldr images with width W, height H, and channel size of 3 (RGB)
20                     exposures(list): list of length N, representing exposures of each images.
21                         Each exposure should correspond to LDR images' exposure value.
22             Return:
23                 (np.ndarray): H x W x 3 shaped numpy array representing HDR image merged u
24                     naive ldr merging implementation.
25                 (np.ndarray): N x H x W x 3 shaped numpy array represending log irradianc
26                     for each exposures
27
28             """
29             N, H, W, C = ldr_images.shape
30             # sanity check
31             assert N == len(exposures)
32
33             # TO DO
34             # Divide each image by exposure time to find irradiance
35             running_sum = np.zeros((H, W, 3))
36             log_irradiances = np.zeros_like(ldr_images)
37
38             for i in range(N):
39                 curr_irradiance = ldr_images[i] / exposures[i]
40                 curr_log_irradiance = np.log(curr_irradiance)
41                 log_irradiances[i] = curr_log_irradiance
42                 running_sum = np.add(running_sum, curr_irradiance)
43
44             # Get average irradiance
45             hdr_image = running_sum / N
46             return hdr_image, log_irradiances
47
48

```

```

1 def display_hdr_image(im_hdr):
2     """

```

```

3     Maps the HDR intensities into a 0 to 1 range and then displays.
4     Three suggestions to try:
5         (1) Take log and then linearly map to 0 to 1 range (see display.py for examp
6             (2) img_out = im_hdr / (1 + im_hdr)
7             (3) HDR display code in a python package
8             ...
9
10    # Log and normalize
11    le = np.log(im_hdr)
12
13    # print('le min max', le.min(), le.max())
14
15    le_min = le[le != -float('inf')].min()
16    le_max = le[le != float('inf')].max()
17    le[le==float('inf')] = le_max
18    le[le==float('-inf')] = le_min
19
20    le = (le - le_min) / (le_max - le_min)
21
22    # Some more scaling magic
23    # img_out = le / (1.0 + le)
24
25
26    # Display image
27    plt.imshow(le)
28

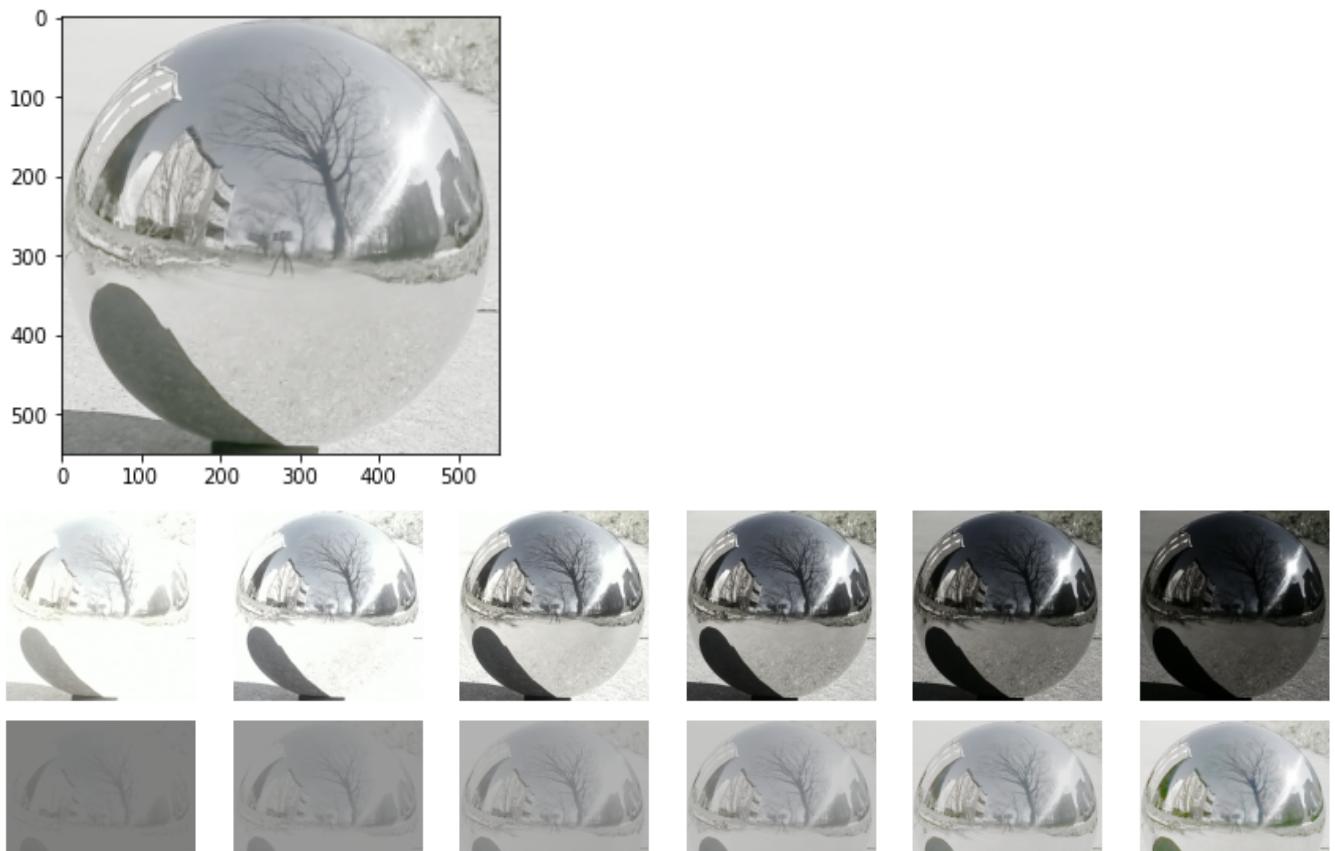
```

```

1
2 # get HDR image, log irradiance
3 naive_hdr_image, naive_log_irradiances = make_hdr_naive(ldr_images, exposure_times
4
5 # write HDR image to directory
6 write_hdr_image(naive_hdr_image, 'images/outputs/naive_hdr.hdr')
7
8 # display HDR image
9 print('HDR Image')
10 display_hdr_image(naive_hdr_image)
11
12 # display original images (code provided in utils.display)
13 display_images_linear_rescale(ldr_images)
14
15 # display log irradiance image (code provided in utils.display)
16 display_images_linear_rescale(naive_log_irradiances)
17

```

```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:40: RuntimeWarning:
HDR Image
```



## ▼ Weighted LDR merging

Compute HDR image as a weighted average of irradiance estimates from LDR images, where weight is based on pixel intensity so that very low/high intensities get less weight

```

1 # this is overcomplicated, but oh well
2 def get_image_weights(img):
3     # weighted_image = np.copy(img)
4
5     # for i in range(img.shape[0]):
6     #     for j in range(img.shape[1]):
7     #         for k in range(img.shape[2]):
8     #             weighted_image[i][j][k] = weighted_image[i][j][k] - 0.5
9
10    # weighted_image = np.abs(weighted_image)
11    # weighted_image = -1 * weighted_image
12
13    # for i in range(img.shape[0]):
14    #     for j in range(img.shape[1]):
15    #         for k in range(img.shape[2]):
16    #             weighted_image[i][j][k] = weighted_image[i][j][k] + 0.5
17
18    # return weighted_image
19    retval = np.double(128.0 - np.abs(np.double(img) - 128.0))
```

```

20     return retval
21
22 w = lambda z: float(128-abs(z-128))
23
24 # clean ldr image by removing values equal to 0 or 255
25 # must pass in rgb values as 0 to 255 scale, with ints not floats
26 # doesn't modify original
27 def clean_ldr_images(ldr_images):
28     N, H, W, C = ldr_images.shape
29     ldr_images_clean = np.int64(255 * np.copy(ldr_images))
30     for n in range(N):
31         for h in range(H):
32             for w in range(W):
33                 for c in range(C):
34                     if (ldr_images_clean[n][h][w][c] == 0):
35                         ldr_images_clean[n][h][w][c] = 1
36                     elif (ldr_images_clean[n][h][w][c] == 255):
37                         ldr_images_clean[n][h][w][c] = 254
38     return ldr_images_clean

1 def make_hdr_weighted(ldr_images: np.ndarray, exposure_times: list) -> (np.ndarray
2     '''
3         Makes HDR image using multiple LDR images, and its corresponding exposure valu
4
5         The steps to implement:
6         1) compute weights for images with based on intensities for each exposures
7             - This can be a binary mask to exclude low / high intensity values
8
9         2) Divide each images by its exposure time.
10            - This will rescale images as if it has been exposed for 1 second.
11
12         3) Return weighted average of above images
13
14
15     Args:
16         ldr_images(np.ndarray): N x H x W x 3 shaped numpy array representing
17             N ldr images with width W, height H, and channel size of 3 (RGB)
18         exposure_times(list): list of length N, representing exposures of each ima
19             Each exposure should correspond to LDR images' exposure value.
20     Return:
21         (np.ndarray): H x W x 3 shaped numpy array representing HDR image merged w
22             under - over exposed regions
23
24     '''
25     N, H, W, C = ldr_images.shape
26     # sanity check
27     assert N == len(exposure_times)
28
29     # TO DO
30
31     # Compute image weights (closer to 0.5 -> 1.0, closer to 0.0 or 1.0 -> 0.0)

```

```

32     ldr_images_clean = clean_ldr_images(ldr_images)
33     image_weights = np.zeros_like(ldr_images)
34
35     for i in range(N):
36         image_weights[i] = get_image_weights(ldr_images_clean[i])
37         image_weights_sum = np.sum(image_weights, axis=0)
38
39     # Divide each image weight by exposure time, then multiply by weights
40     running_sum = np.zeros((H, W, 3))
41     for i in range(N):
42         curr_irradiance = ldr_images[i] / exposure_times[i]
43         curr_irradiance_weights = np.multiply(curr_irradiance, image_weights[i])
44         running_sum = np.add(running_sum, curr_irradiance_weights)
45
46     # Return weighted image average
47     # (A * wA + B * wB + ...) / (wA + wB + ...)
48     weighted_average = np.divide(running_sum, image_weights_sum)
49     return weighted_average
50
51

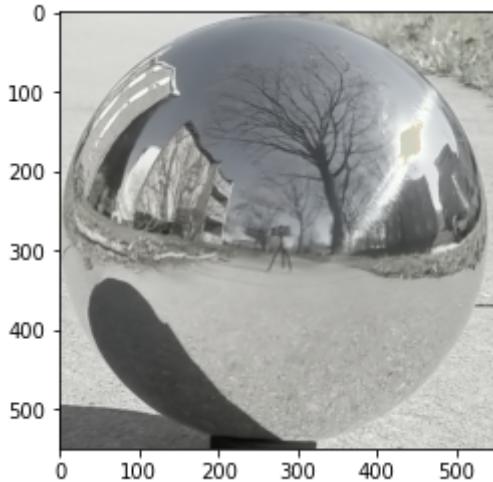
```

```

1 # get HDR image, log irradiance
2 weighted_hdr_image = make_hdr_weighted(ldr_images, exposure_times)
3
4 # write HDR image to directory
5 write_hdr_image(weighted_hdr_image, 'images/outputs/weighted_hdr.hdr')
6
7 # display HDR image
8 print("Weighted HDR image:")
9 display_hdr_image(weighted_hdr_image)
10

```

Weighted HDR image:

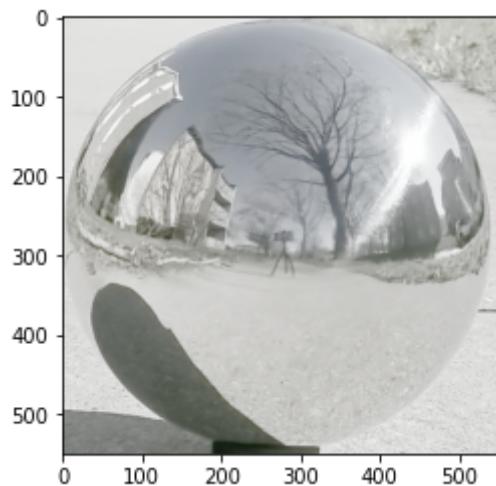


```

1 print('Naive HDR')
2 display_hdr_image(naive_hdr_image)

```

Naive HDR



Display of difference between naive and weighted for your own inspection

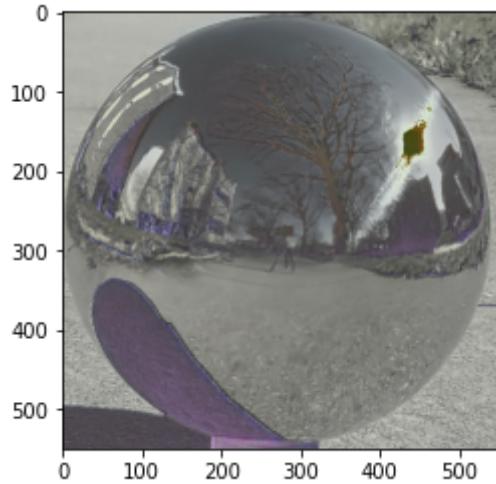
Where does the weighting make a big difference increasing or decreasing the irradiance estimate?  
Think about why.

```

1 # display difference between naive and weighted
2
3 log_diff_im = np.log(weighted_hdr_image)-np.log(naive_hdr_image)
4 print('Min ratio = ', np.exp(log_diff_im).min(), ' Max ratio = ', np.exp(log_diff_im).max())
5 plt.figure()
6 plt.imshow(rescale_images_linear(log_diff_im))

```

Min ratio = 0.665383278598632 Max ratio = 2.7496184788243125  
<matplotlib.image.AxesImage at 0x7f71a4107110>



## ▼ LDR merging with camera response function estimation

Compute HDR after calibrating the photometric responses to obtain more accurate irradiance estimates from each image

## Some suggestions on using gsolve:

- When providing input to gsolve, don't use all available pixels, otherwise you will likely run out of memory / have very slow run times. To overcome, just randomly sample a set of pixels (1000 or so can suffice), but make sure all pixel locations are the same for each exposure.
- The weighting function w should be implemented using Eq. 4 from the paper (this is the same function that can be used for the previous LDR merging method).
- Try different lambda values for recovering g. Try lambda=1 initially, then solve for g and plot it. It should be smooth and continuously increasing. If lambda is too small, g will be bumpy.
- Refer to Eq. 6 in the paper for using g and combining all of your exposures into a final image. Note that this produces log irradiance values, so make sure to exponentiate the result and save irradiance in linear scale.

```

1 # Sample random points from the mirror ball region
2 # Assumes mirror ball is perfectly inscribed inside square image
3 # Image has dimensions H x W x 3 (we must take all channels for a pixel, of course
4 # P is number of pixels to sample
5 # Returns list of pixel coordinates (tuples of length 2)
6 def get_random_pixels(H, W, P):
7     all_pixels = []
8     for i in range(H):
9         for j in range(W):
10            all_pixels.append((i, j))
11    return random.sample(all_pixels, P)
12
13 # return 3 x N x P array to feed into g function
14 # note: only N x P will be fed in (one layer at a time)
15 def get_NP_array(ldr_images, pixels):
16     N = len(ldr_images)
17     P = len(pixels)
18     NP_array = np.zeros((3, N, P))
19     for i in range(3):
20         for j in range(N):
21             for k in range(P):
22                 NP_array[i][j][k] = ldr_images[j][pixels[k][0]][pixels[k][1]][i]
23     return np.int64(NP_array)
24

```

```

1 def make_hdr_estimation(ldr_images: np.ndarray, exposure_times: list, lm) -> (np.nd
2     ''
3     Makes HDR image using multiple LDR images, and its corresponding exposure value
4     Please refer to problem notebook for how to do it.
5
6     **IMPORTANT**
7     The gsolve operations should be ran with:
8         Z: int64 array of shape N x P, where N = number of images, P = number of p

```

```

9         B: float32 array of shape N, log shutter times
10        l: lambda; float to control amount of smoothing
11        w: function that maps from float intensity to weight
12 The steps to implement:
13 1) Create random points to sample (from mirror ball region)
14 2) For each exposures, compute g values using samples
15 3) Recover HDR image using g values
16
17
18 Args:
19     ldr_images(np.ndarray): N x H x W x 3 shaped numpy array representing
20         N ldr images with width W, height H, and channel size of 3 (RGB)
21     exposures(list): list of length N, representing exposures of each images.
22         Each exposure should correspond to LDR images' exposure value.
23     lm (scalar): the smoothing parameter
24 Return:
25     (np.ndarray): H x W x 3 shaped numpy array representing HDR image merged u
26         gsolve
27     (np.ndarray): N x H x W x 3 shaped numpy array represending log irradiance
28         for each exposures
29     (np.ndarray): 3 x 256 shaped numpy array represending g values of each pix
30         at each channels (used for plotting)
31     ...
32     N, H, W, C = ldr_images.shape
33     # sanity check
34     assert N == len(exposure_times)
35
36     # TO DO: implement HDR estimation using gsolve
37     # gsolve(Z, B, l, w) -> g, LE
38
39     # We will use 10000 pixels for gsolve
40     P = 10000
41     sample_pixels = get_random_pixels(H, W, P)
42     log_exposure_times = np.log(exposure_times)
43     ldr_images_clean = clean_ldr_images(ldr_images)
44     NP_array = get_NP_array(ldr_images_clean, sample_pixels)
45
46     hdr_image = np.ones((H, W, C))
47     log_irradiances = np.zeros_like(ldr_images)
48     gs = np.ones((C, 256))
49     LEs = []
50
51     # Get g-values
52     # For each image, there will be a g-value for EACH channel
53     weights = np.arange(256)
54     weights = get_image_weights(weights)
55     for i in range(C):
56         g, LE = gsolve(NP_array[i], log_exposure_times, lm, weights)
57         gs[i] = g
58         LEs.append(LE)
59

```

```

60     # reconstruct HDR image
61     # use equation 6 from paper
62     log_radiance = np.zeros((H, W, C))
63     for c in range(C):
64         # current image-channel
65         for h in range(H):
66             for w in range(W):
67                 uppersum = 0
68                 lowersum = 0
69                 for n in range(N):
70                     curr_pixel_value = ldr_images_clean[n, h, w, c]
71                     uppersum += get_image_weights(curr_pixel_value) * (gs[c][curr_pixel_va
72                     lowersum += get_image_weights(curr_pixel_value)
73                     log_radiance[h][w][c] = float(uppersum) / float(lowersum)
74
75     hdr_image = np.exp(log_radiance)
76
77     # calculate log irradiances
78     for i in range(N):
79         curr_irradiance = ldr_images[i] / exposure_times[i]
80         curr_log_irradiance = np.log(curr_irradiance)
81         log_irradiances[i] = curr_log_irradiance
82         # log_irradiances[i] = log_radiance
83
84     return hdr_image, log_irradiances, gs
85

```

```

1 lm = 10
2 # get HDR image, log irradiance
3 calib_hdr_image, calib_log_irradiances, g = make_hdr_estimation(ldr_images, exposu
4 fig = plt.figure(figsize=(6, 6))
5 plt.xlabel('intensity')
6 plt.ylabel('g-value')
7 plt.title('intensity vs. g-value')
8 plt.scatter(np.arange(256), g[0],
9             linewidths=1, alpha=.7,
10            edgecolor='k',
11            c='red', label='red')
12 plt.scatter(np.arange(256), g[1],
13             linewidths=1, alpha=.7,
14            edgecolor='k',
15            c='green', label='green')
16 plt.scatter(np.arange(256), g[2],
17             linewidths=1, alpha=.7,
18            edgecolor='k',
19            c='blue', label='blue')
20 plt.legend(title='channel')
21 plt.show()
22
23 # write HDR image to directory

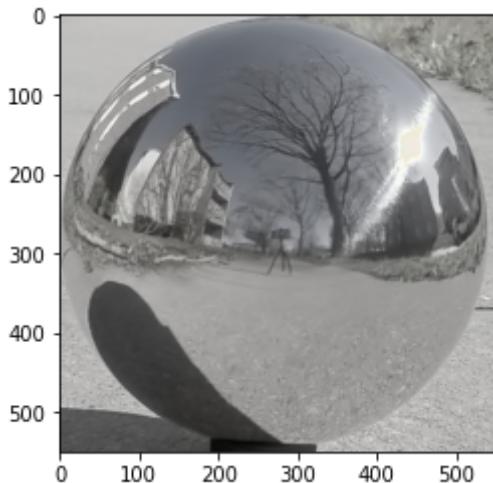
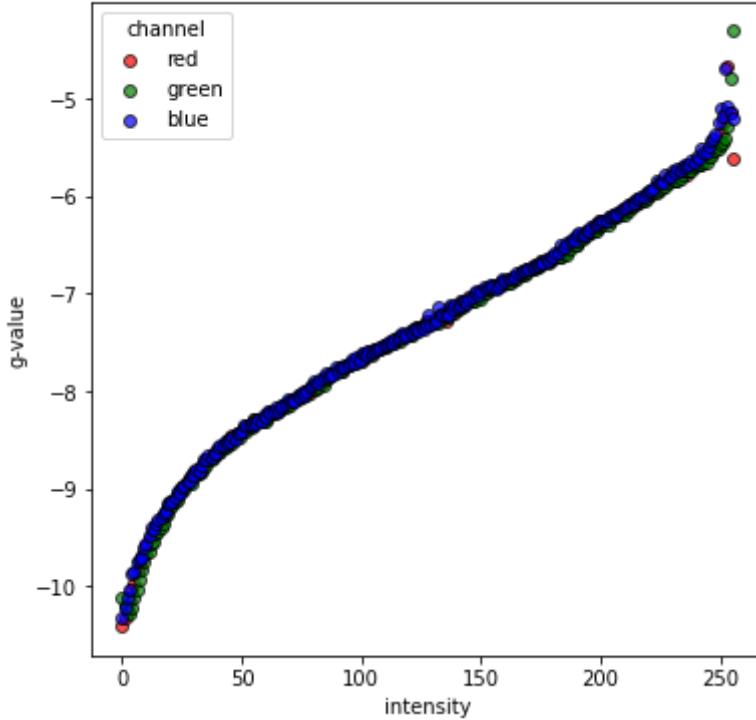
```

```

24 write_hdr_image(calib_hdr_image, 'images/outputs/calib_hdr.hdr')
25
26 # display HDR image
27 display_hdr_image(calib_hdr_image)
28

```

/usr/local/lib/python3.7/dist-packages/ipykernel\_launcher.py:80: RuntimeWarning:  
intensity vs. g-value



The following code displays your results. You can copy the resulting images and plots directly into your report where appropriate.

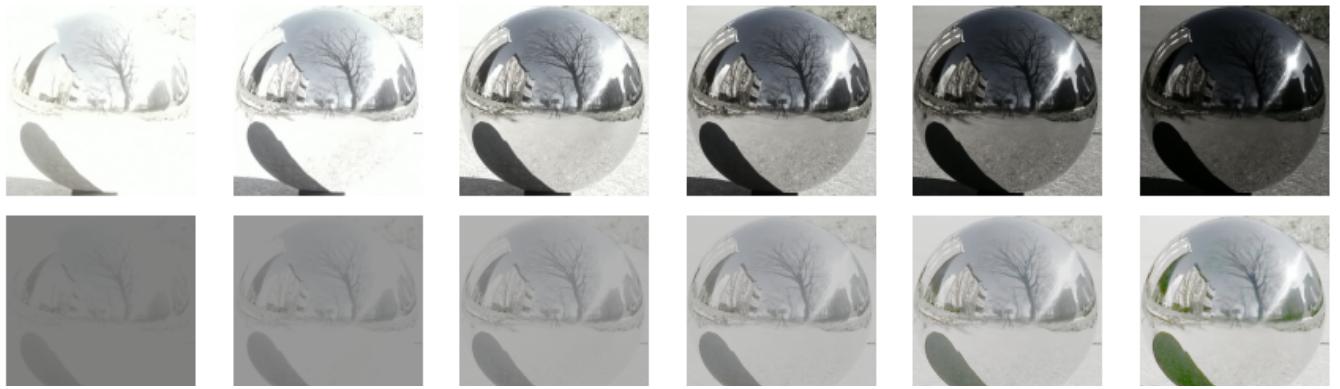
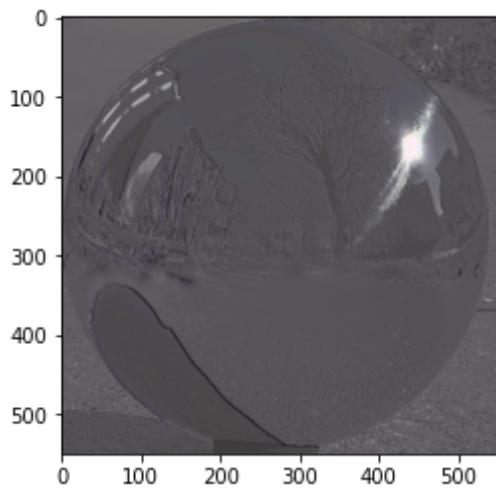
```

1 # display difference between calibrated and weighted
2 log_diff_im = np.log(calib_hdr_image/calib_hdr_image.mean())-np.log(weighted_hdr_i
3 print('Min ratio = ', np.exp(log_diff_im).min(), ' Max ratio = ', np.exp(log_diff
4 plt.figure()

```

```
5 plt.imshow(rescale_images_linear(log_diff_im))
6
7 # display original images (code provided in utils.display)
8 display_images_linear_rescale(ldr_images)
9
10 # display log irradiance image (code provided in utils.display)
11 display_images_linear_rescale(calib_log_irradiances)
12
13 # plot g vs intensity, and then plot intensity vs g
14 N, NG = g.shape
15 labels = [ 'R', 'G', 'B' ]
16 plt.figure()
17 for n in range(N):
18     plt.plot(g[n], range(NG), label=labels[n])
19 plt.gca().legend(( 'R', 'G', 'B' ))
20
21 plt.figure()
22 for n in range(N):
23     plt.plot(range(NG), g[n], label=labels[n])
24 plt.gca().legend(( 'R', 'G', 'B' ))
```

```
Min ratio = 0.48564966732012194 Max ratio = 3.218209877178028
<matplotlib.legend.Legend at 0x7f71a1cbf310>
```



```
1 def weighted_log_error(ldr_images, hdr_image, log_irradiances):
2     # computes weighted RMS error of log irradiances for each image compared to final
3     N, H, W, C = ldr_images.shape
4     w = 1-abs(ldr_images - 0.5)*2
5     err = 0
6     for n in np.arange(N):
7         err += np.sqrt(np.multiply(w[n], (log_irradiances[n]-np.log(hdr_image))**2).sum())
8     return err
9
10
11 # compare solutions
12 err = weighted_log_error(ldr_images, naive_hdr_image, naive_log_irradiances)
13 print('naive:\tlog range = ', round(np.log(naive_hdr_image).max() - np.log(naive_hdr_image).min(), 2))
14 err = weighted_log_error(ldr_images, weighted_hdr_image, naive_log_irradiances)
15 print('weighted:\tlog range = ', round(np.log(weighted_hdr_image).max() - np.log(weighted_hdr_image).min(), 2))
16 err = weighted_log_error(ldr_images, calib_hdr_image, calib_log_irradiances)
17 print('calibrated:\tlog range = ', round(np.log(calib_hdr_image).max() - np.log(calib_hdr_image).min(), 2))
18
19 # display log hdr images (code provided in utils.display)
20 display_images_linear_rescale(np.log(np.stack((naive_hdr_image/naive_hdr_image.mean(), weighted_hdr_image/weighted_hdr_image.mean(), calib_hdr_image/calib_hdr_image.mean()), axis=2)))
```

naive:	log range = 4.216	avg RMS error = 0.454
weighted:	log range = 4.466	avg RMS error = 0.407
calibrated:	log range = 5.396	avg RMS error = 6.419



## ▼ Panoramic transformations

Compute the equirectangular image from the mirrorball image

```

1 from re import U
2 def panoramic_transform(hdr_image):
3     '''
4         Given HDR mirror ball image,
5
6         Expects mirror ball image to have center of the ball at center of the image, a
7         width and height of the image to be equal.
8
9         Steps to implement:
10        1) Compute N image of normal vectors of mirror ball
11        2) Compute R image of reflection vectors of mirror ball
12        3) Map reflection vectors into spherical coordinates
13        4) Interpolate spherical coordinate values into equirectangular grid.
14
15        Steps 3 and 4 are implemented for you with get_equirectangular_image
16
17        '''
18        H, W, C = hdr_image.shape
19        assert H == W
20        assert C == 3
21
22        # TO DO: compute N and R
23
24        # V is constant looking into the image
25        V = np.float32(np.array([0,0,-1]))
26
27        # N must be calculated based on reflection off sphere
28        center = H / 2.0
29        N = np.float32(np.zeros((H, W, C)))
30        for h in range(H):
31            for w in range(W):
32                N[h][w][0] = (w - center) / (W / 2.0)

```

```

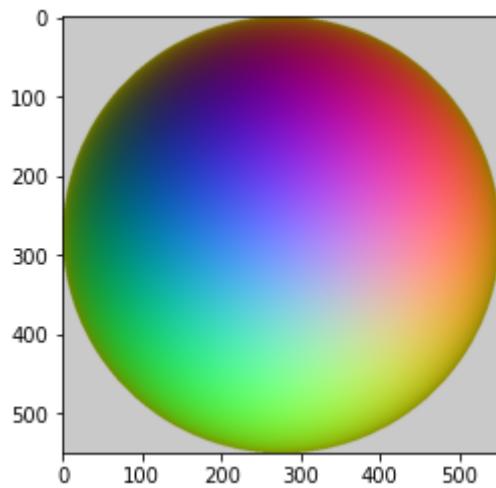
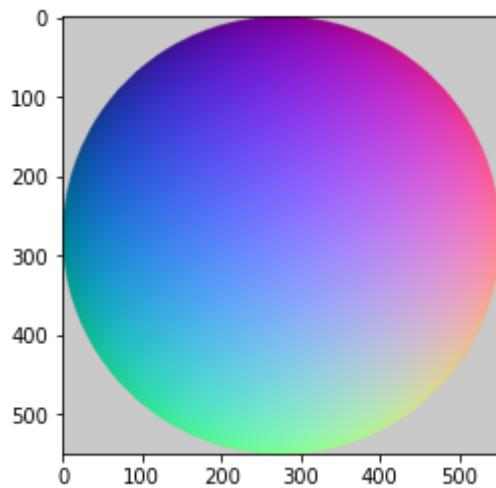
33     N[h][w][1] = (h - center) / (H / 2.0)
34     # must check if outside sphere of influence (no pun intended)
35     if (N[h][w][0]**2 + N[h][w][1]**2 >= 1.0):
36         N[h][w][0] = 1.0
37         N[h][w][1] = 1.0
38         N[h][w][2] = 1.0
39     else:
40         N[h][w][2] = np.sqrt(1.0 - N[h][w][0]**2 - N[h][w][1]**2)
41         # normalize the normals
42     N[h][w] = np.float32(N[h][w] / np.linalg.norm(N[h][w]))
43
44     # R = V - 2 * dot(V,N) * N
45     R = np.float32(np.zeros((H, W, C)))
46     for h in range(H):
47         for w in range(W):
48             Nx = (np.float32(w) - center) / (W / 2.0)
49             Ny = (np.float32(h) - center) / (H / 2.0)
50             # must check if outside sphere of influence (no pun intended)
51             if (Nx**2 + Ny**2 >= 1.0):
52                 R[h][w][0] = 1.0
53                 R[h][w][1] = 1.0
54                 R[h][w][2] = 1.0
55             else:
56                 R[h][w] = np.float32(V - 2.0 * np.dot(V, N[h][w]) * N[h][w])
57             R[h][w] = np.float32(R[h][w] / np.linalg.norm(R[h][w]))
58
59     plt.imshow((N+1)/2)
60     plt.show()
61     plt.imshow((R+1)/2)
62     plt.show()
63
64     equirectangular_image = get_equirectangular_image(R, hdr_image)
65     return equirectangular_image, N, R

```

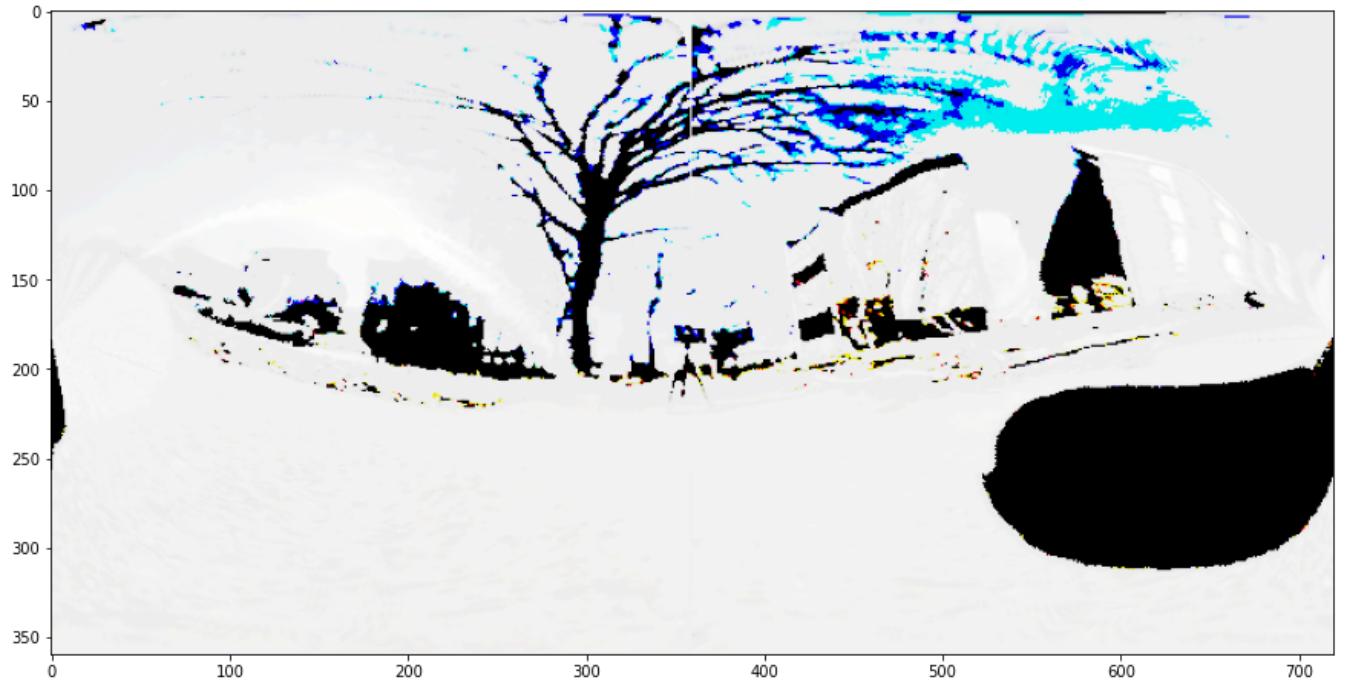
```

1 hdr_mirrorball_image = read_hdr_image('images/outputs/calib_hdr.hdr')
2 eq_image, N, R = panoramic_transform(hdr_mirrorball_image)
3 eq_image = np.abs(eq_image)
4
5 write_hdr_image(eq_image, 'images/outputs/equirectangular.hdr')
6
7 plt.figure(figsize=(15,15))
8 display_hdr_image(eq_image)
9

```



```
/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:11: RuntimeWarning:  
# This is added back by InteractiveShellApp.init_path()
```



## ▼ Rendering synthetic objects into photographs

Use Blender to render the scene with and with objects and obtain the mask image. The code below should then load the images and create the final composite.

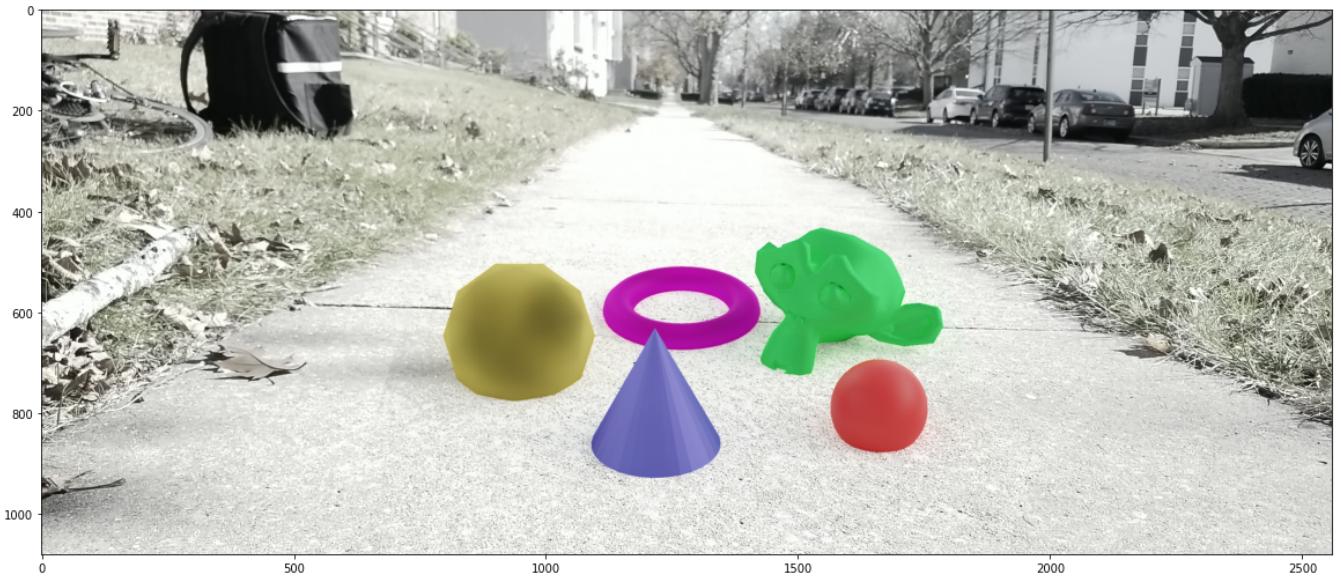
Note: images must be uploaded manually to the local images directory.

```
1 # Read the images that you produced using Blender.  Modify names as needed.
2 O = read_image('images/p41_rendered.png')
3 E = read_image('images/p41_empty_rendered.png')
4 M = read_image('images/p41_mask_rendered.png')
5 M = M > 0.5
6 I = read_image('images/p41_background.jpg')
7 I = cv2.resize(I, (M.shape[1], M.shape[0]))

1 # TO DO: compute final composite
2 c = 1
3 M_actual = 1.0 - M
4 result = np.multiply(M_actual, O) + np.multiply(1.0-M_actual, I) + np.multiply(1.0
5
6 plt.figure(figsize=(20,20))
7 plt.imshow(result)
8 plt.show()
9
10 write_image(result, 'images/outputs/final_composite_1.png')
11
12
13
14
15
16
17
18
19
20
21
22
23
24
25
26
27
28
29
30
31
```



WARNING:matplotlib.image:Clipping input data to the valid range for imshow with



## ▼ Bells & Whistles (Extra Points)

### Additional Image-Based Lighting Result

Different objects using same light map. Features translucent skyscraper, glass golf trophy, Mandalorian helmet, 50s-era electric train with pantographs, monkey (default Blender object), and reflective teapot.

```

1 # Read the images that you produced using Blender.  Modify names as needed.
2 O = read_image('images/p42_rendered.png')
3 E = read_image('images/p42_empty_rendered.png')
4 M = read_image('images/p42_mask_rendered.png')
5 M = M > 0.5
6 I = read_image('images/p42_background.jpg')
7 I = cv2.resize(I, (M.shape[1], M.shape[0]))

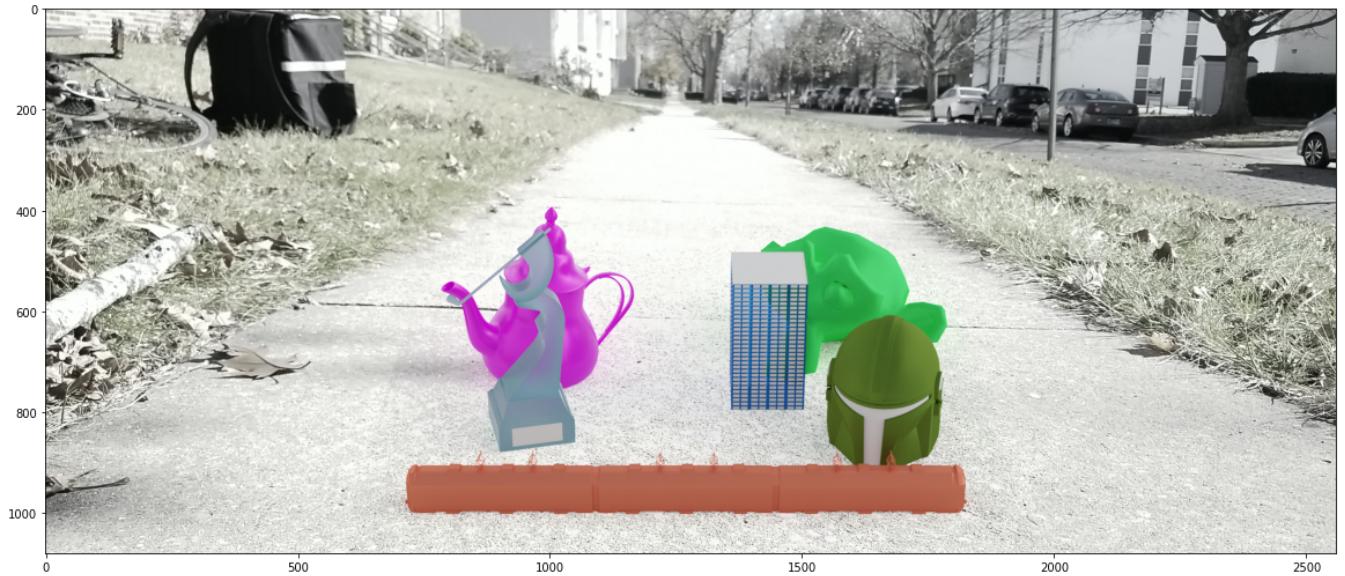

1 # TO DO: compute final composite
2 c = 1
3 M_actual = 1.0 - M
4 result = np.multiply(M_actual, O) + np.multiply(1.0-M_actual, I) + np.multiply(1.0
-
```

```

5
6 plt.figure(figsize=(20,20))
7 plt.imshow(result)
8 plt.show()
9
10 write_image(result, 'images/outputs/final_composite_2.png')

```

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with



Same objects but different light map. I used my apartment as the setting, with the ball placed upon a desk. Lighting is radically different (a lot dimmer).

```

1 # Read the images that you produced using Blender.  Modify names as needed.
2 O = read_image('images/p43_rendered.png')
3 E = read_image('images/p43_empty_rendered.png')
4 M = read_image('images/p43_mask_rendered.png')
5 M = M > 0.5
6 I = read_image('images/p43_background.jpg')
7 I = cv2.resize(I, (M.shape[1], M.shape[0]))

1 # TO DO: compute final composite
2 c = 1

```

```
3 M_actual = 1.0 - M
4 result = np.multiply(M_actual, O) + np.multiply(1.0-M_actual, I) + np.multiply(1.0
5
6 plt.figure(figsize=(20,20))
7 plt.imshow(result)
8 plt.show()
9
10 write_image(result, 'images/outputs/final_composite_3.png')
```

WARNING:matplotlib.image:Clipping input data to the valid range for imshow with



Colab paid products - Cancel contracts here

✓ 2s completed at 9:33 PM

